

Applied Text Mining in Python

Naïve Bayes Classifier

Case study: Classifying text search queries

- Suppose you are interested in classifying search queries in three classes: **Entertainment**, **Computer Science**, **Zoology**
- Most common class of the three is **Entertainment**.

Case study: Classifying text search queries

- Suppose the query is “Python”
 - Python, the snake (Zoology)
 - Python, the programming language (Computer Science)
 - Python, as in Monty Python (Entertainment)
- Most common class, given “Python”, is Zoology.

Case study: Classifying text search queries

- Suppose the query is “Python download”
- Most probable class, given “Python download”, is Computer Science.

Probabilistic Model

- Update the likelihood of the class given new information
- **Prior Probability:** $\Pr(y = \text{Entertainment})$, $\Pr(y = \text{CS})$, $\Pr(y = \text{Zoology})$
- **Posterior probability:** $\Pr(y = \text{Entertainment} | x = \text{"Python"})$

Bayes' Rule

- Posterior probability =
$$\frac{\text{Prior probability} \times \text{Likelihood}}{\text{Evidence}}$$
- $$\Pr(y | X) = \frac{\Pr(y) \times \Pr(X | y)}{\Pr(X)}$$

Naïve Bayes Classification

- $\Pr(y=\text{CS} | \text{"Python"}) = \frac{\Pr(y=\text{CS}) \times \Pr(\text{"Python"} | y=\text{CS})}{\Pr(\text{"Python"})}$
- $\Pr(y=\text{Zoology} | \text{"Python"}) = \frac{\Pr(y=\text{Zoology}) \times \Pr(\text{"Python"} | y=\text{Zoology})}{\Pr(\text{"Python"})}$
- $\Pr(y=\text{CS} | \text{"Python"}) > \Pr(y=\text{Zoology} | \text{"Python"}) \Rightarrow$

Naïve Bayes Classification

- $\Pr(y | X) = \frac{\Pr(y) \times \Pr(X | y)}{\Pr(X)}$

$$y^* = \underset{y}{\operatorname{argmax}} \Pr(y | X) = \underset{y}{\operatorname{argmax}} \Pr(y) \times \Pr(X | y)$$

- **Naïve assumption:** Given the class label, features are assumed to be independent of each other

$$y^* = \underset{y}{\operatorname{argmax}} \Pr(y | X) = \underset{y}{\operatorname{argmax}} \Pr(y) \times \prod_{i=1}^n \Pr(x_i | y)$$

Naïve Bayes Classifier

$$y^* = \operatorname{argmax}_y \Pr(y \mid X) = \operatorname{argmax}_y \Pr(y) \times \prod_{i=1}^n \Pr(x_i \mid y)$$

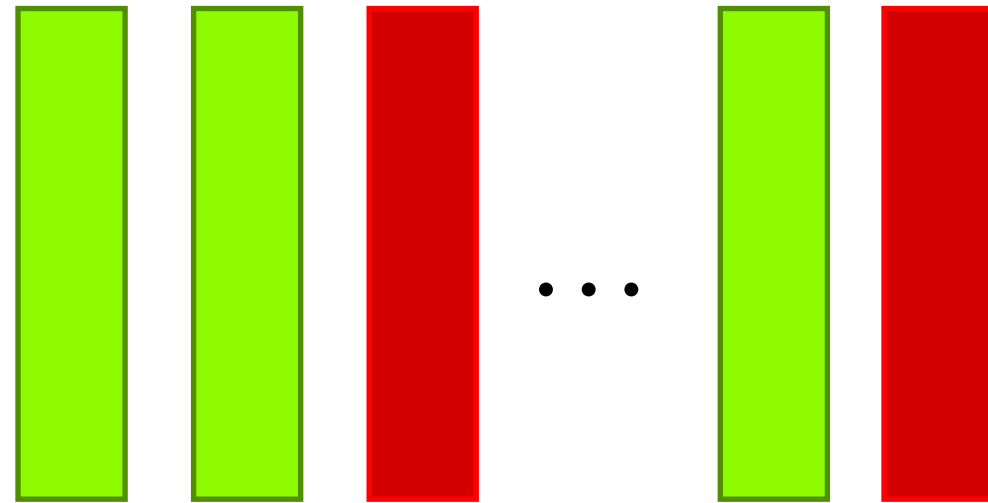
Query: “Python download”

$$y^* = \operatorname{argmax}_y \Pr(y) \times \Pr(\text{“Python”} \mid y) \times \Pr(\text{“download”} \mid y)$$

Naïve Bayes: What are the parameters?

- Prior probabilities: $\Pr(y)$ for all y in Y
- Likelihood: $\Pr(x_i | y)$ for all features x_i and labels y in Y
- If there are 3 classes ($|Y| = 3$) and 100 features in X , how many parameters does naïve Bayes models have?

Naïve Bayes: Learning parameters

- Prior probabilities: $\Pr(y)$ for all y in Y
- Remember training data?
- Count the number of instances in each class
- If there are N instances in all, and n out of those are labeled as class $y \Rightarrow \Pr$

Naïve Bayes: Learning parameters

- Likelihood: $\Pr(x_i | y)$ for all features x_i and labels y in Y
- Count how many times feature x_i appears in instances labeled as class y
- If there are p instances of class y , and x_i appears in k of those, $\Pr(x_i | y) = k / p$

Naïve Bayes: Smoothing

- What happens if $\Pr(\mathbf{x}_i \mid \mathbf{y}) = 0$?
 - Feature \mathbf{x}_i never occurs in documents labeled \mathbf{y}
 - But then, the posterior probability $\Pr(\mathbf{y} \mid \mathbf{x}_i)$ will be 0!!
- Instead, smooth the parameters
- **Laplace smoothing** or **Additive smoothing**: Add a dummy count
- $\Pr(\mathbf{x}_i \mid \mathbf{y}) = (k+1) / (p+n)$; where n is number of features

Take Home Concepts

- **Naïve Bayes is a probabilistic model**
- **Naïve, because it assumes features are independent of each other, given the class label**
- **For text classification problems, naïve Bayes models typically provide very strong baselines**
- **Simple model, easy to learn parameters**